



October 4-6, 2017 | Vancouver, BC

Node.js With Steroids

Make Better Node.js Application with Native Add-Ons

Nicola Del Gobbo, Developer, *Packly*



October 4-6, 2017 | Vancouver, BC

Who is Nicola Del Gobbo?

Developer



Contribute to



open source




What is Node.js Native Addon?

In a very simple way Native Addons could be considered as
C / C++ code called from JavaScript



Native environment



JavaScript environment





October 4-6, 2017 | Vancouver, BC

From Node.js documentation

*Node.js Addons are **dynamically-linked shared objects**, written in **C++**, that can be loaded into Node.js using the **require()** function, and used just as if they were an ordinary Node.js module.*

*They are used primarily to provide an **interface** between **JavaScript** running in Node.js and **C/C++** libraries.*



How is it possible?

All the **magic** is behind the **Node.js** architecture

Node.js API

Node.js Bindings

C / C++ Addons



c-ares

HTTP
parser

Open
SSL

zlib

nghttp2

Why should you implement new Native Add-on?

Performance

In general C / C++ code performs better than JavaScript code, but it's really true for **CPU bounds** operations

- **Image** processing (in average 6 times faster)
- **Video** processing
- **CRC** cyclic redundancy check (in average 125 time faster)
- **Compression**
- **Scientific calculus**
- **Algorithms that execute CPU heavy tasks**




Why should you implement new Native Add-on?

Integrate legacy application

You have the source code of an **old C / C++** application and want to expose something of its functionalities through the new Node.js application

```
#include <stdio.h>
int main(void)
{
    printf("Hello World!\n");
    return 0;
}
```

C/C++



Why should you implement new Native Add-on?

You don't find what fits your specific needs on **npm**

Sometimes completely reimplementing the module in **JavaScript** is not the right solution

Think at libraries like:

- **ImageMagick**
- **Ghostscript**
- **FFmpeg**
- **TensorFlow**

You have to concentrate on business logic of your application and not on developing a single module



Why should you implement new Native Add-on?

Desktop application

You want to use the local underlying **CPU** or **GPU** resources to execute the most complex tasks

- Better performance on the **local** side
- Better performance on the **server**

Es. **Electron** or **NW.js**



Why should you implement new Native Add-on?

Better error handling

Even if you will use **an old C / C++** library you will get an error code that explains the error that just happened.

In **new C / C++** library you have the **exception**

You don't have to parse some string to identify if there was an error and what kind of it



Problems

Fragmentation API

The API to implement Native Add-Ons have been changed across different version of Node.js

Most of the changes were on **V8 API** and **ObjectWrap API**



0.8 - 0.10.x - 0.12.x - 1.x - 2.x - 3.x - 4.x - 5.x - 6.x - 7.x - 8.x

For more info <http://v8docs.nodesource.com/>

Problems

Fragmentation API

Need an adapter to stay compatible across different version of Node.js

- **NAN** - Native Abstraction for Node.js
 - **API compatibility**
 - Strong bonded with **V8 API**
 - You have to recompile your native add-ons switching to different version of Node.js



Problems

Write portable C / C++ code

At one point your native code **must compile** on different:

ARCHITECTURE

PLATFORM

COMPILER



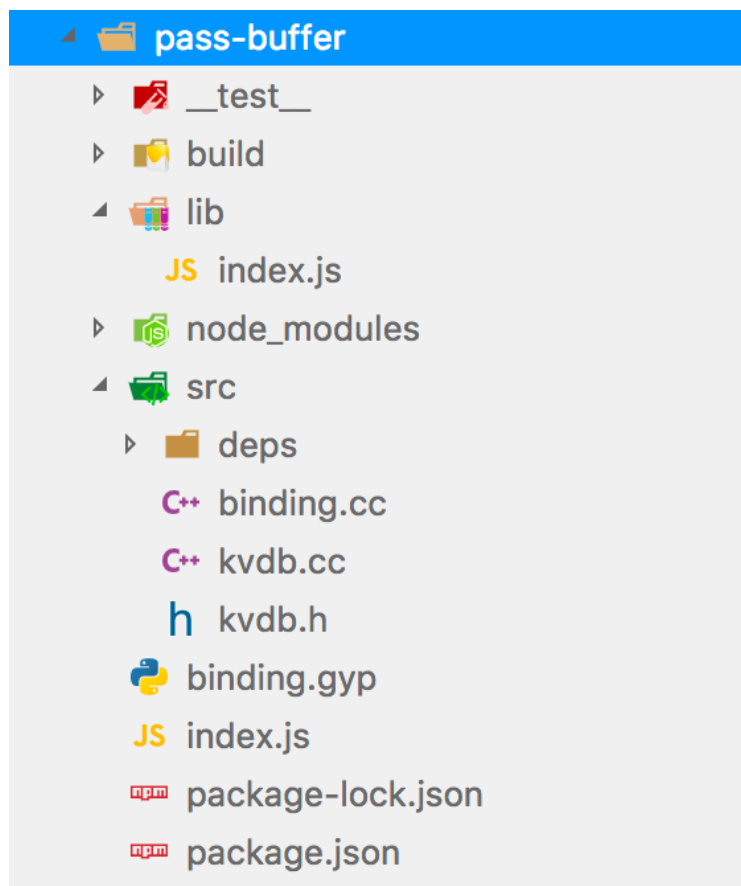
Problems

Documentation

- **C / C++** libraries that you are integrating are not well documented
- There are good **references** but not so much practical guide focusing on concepts about Native Add-Ons



How to organize your native Addon project



src C/C++ code and maybe dependencies

binding.gyp building configurations

lib JavaScript code

package.json

Example: the JavaScript part

```
'use strict';

var binary = require('node-pre-gyp');
var path = require('path');
var binding_path = binary.find(path.resolve(path.join(__dirname, './package.json')));
```

Import the Addon

```
var bindings = require(binding_path);

var crypto = require('crypto');

var promises = require('./lib/promises');
```

```
/// generate a salt (sync)
/// @param {Number} [rounds] number of rounds (default 10)
/// @return {String} salt
module.exports.genSaltSync = function genSaltSync(rounds) {
  // default 10 rounds
  if (!rounds) {
    rounds = 10;
  } else if (typeof rounds !== 'number') {
    throw new Error('rounds must be a number');
  }
}
```

Call the Addon's function

```
return bindings.gen_salt_sync(rounds, crypto.randomBytes(16));
};
```

```
{
  "name": "conf-ni-2017-4",
  "description": "This is a project used as example for Node Interactive 2017",
  "version": "0.0.1",
  "private": true,
  "author": "Nicola Del Gobbo <nicoladelgobbo@gmail.com>",
  "keywords": [
    "c++",
    "native",
    "module",
    "nodejs",
    "addon",
    "v8"
  ],
  "license": "Apache-2.0",
  "scripts": {
    "test": "node node_modules/jasmine/bin/jasmine.js JASMINE_CONFIG_PATH=__test__/jasmine.json",
    "build": "node-gyp rebuild"
  },
  "engines": {
    "node": ">= 0.10.0"
  },
  "repository": {
    "type": "git",
    "url": "https://github.com/NickNaso/conf-ni-2017.git"
  },
  "bugs": {
    "url": "https://github.com/NickNaso/conf-ni-2017/issues"
  },
  "main": "index.js",
  "devDependencies": {
    "jasmine": "^2.8.0"
  },
  "gypfile": true,
  "dependencies": {
    "bindings": "^1.3.0",
    "nan": "^2.7.0"
  }
}
```


Example: the C / C++ part

```
NAN_METHOD(GenerateSaltSync) {  
    Nan::HandleScope scope;
```

```
    if (info.Length() < 2) {  
        Nan::ThrowTypeError("2 arguments expected");  
        return;  
    }
```

Validate your input

```
    if (!Buffer::HasInstance(info[1]) || Buffer::Length(info[1].As<Object>()) != 16) {  
        Nan::ThrowTypeError("Second argument must be a 16 byte Buffer");  
        return;  
    }
```

Pass input values from JS context to C++

```
    const int32_t rounds = Nan::To<int32_t>(info[0]).FromMaybe(0);  
    u_int8_t* seed = (u_int8_t*)Buffer::Data(info[1].As<Object>());
```

Execute elaboration on the C++ environment

```
    char salt[_SALT_LEN];  
    bcrypt_gensalt(rounds, seed, salt);
```

Return value to JS context

```
    info.GetReturnValue().Set(Nan::Encode(salt, strlen(salt), Nan::BINARY));
```

```
}
```

```
NAN_MODULE_INIT(init) {
```

Export function

```
Nan::Export(target, "gen_salt_sync", GenerateSaltSync);
```

```
Nan::Export(target, "encrypt_sync", EncryptSync);  
Nan::Export(target, "compare_sync", CompareSync);  
Nan::Export(target, "get_rounds", GetRounds);  
Nan::Export(target, "gen_salt", GenerateSalt);  
Nan::Export(target, "encrypt", Encrypt);  
Nan::Export(target, "compare", Compare);
```

```
};
```

Registration and initialization

```
NODE_MODULE(bcrypt_lib, init);
```

Example: the binding.gyp

```
{
  'targets': [
    {
      'target_name': 'bcrypt_lib',
      'sources': [
        'src/blowfish.cc',
        'src/bcrypt.cc',
        'src/bcrypt_node.cc'
      ],
      'include_dirs': [
        "<!(node -e \"require('nan')\")"
      ],
      'conditions': [
        [ 'OS=="win"', {
          'defines': [
            'uint=unsigned int',
          ],
        } ],
      ],
    },
  ],
  {
    "target_name": "action_after_build",
    "type": "none",
    "dependencies": [ "<(module_name)" ],
    "copies": [
      {
        "files": [ "<(PRODUCT_DIR)/<(module_name).node" ],
        "destination": "<(module_path)"
      }
    ]
  }
]
```

Name used to register the Addon

Sometimes finding the right **settings** for binding.gyp is not easy


Refer to **GYP** documentation

<https://gyp.gsrc.io/docs/UserDocumentation.md>

More useful examples to take inspiration from other developers

[https://github.com/nodejs/node-gyp/wiki/"binding.gyp"-files-out-in-the-wild](https://github.com/nodejs/node-gyp/wiki/)

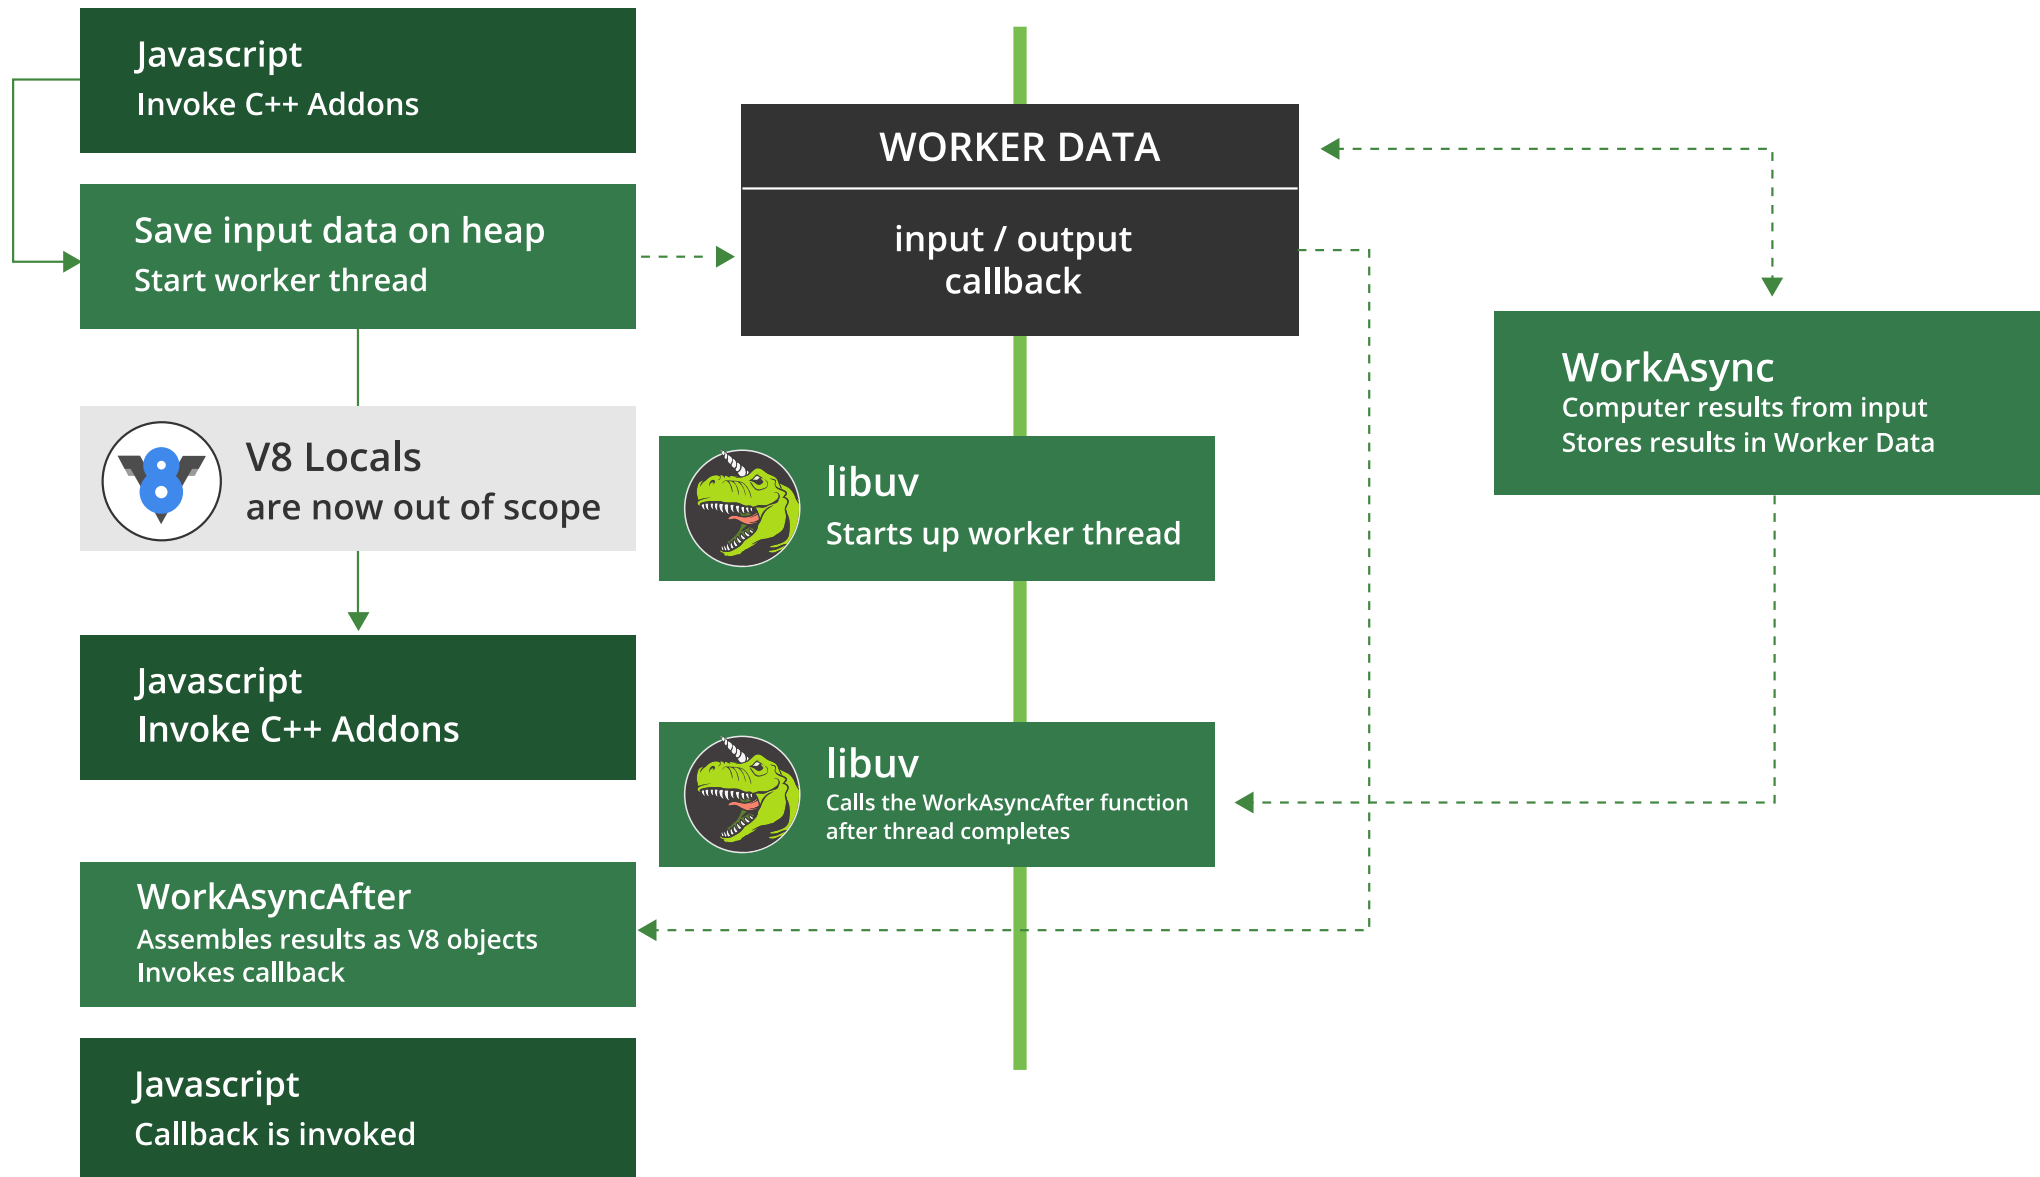
Which kind of Addon should you implement?



TIME: Amount of processing time spent on the **C++** code

DATA: Amount of data flowing between **C++** and **JavaScript**

The async way



Async example

```
class GhostscriptWorker : public AsyncWorker
{
public:
    GhostscriptWorker(Callback *callback, string RAWcmd)
        : AsyncWorker(callback), RAWcmd(RAWcmd), res(0) {}
    ~GhostscriptWorker() {}
```

```
void Execute()
{
    res = 0;
    vector<string> explodedCmd;
    istringstream iss(RAWcmd);
    for (string RAWcmd; iss >> RAWcmd;)
        explodedCmd.push_back(RAWcmd);
    int gsargc = static_cast<int>(explodedCmd.size());
    char **gsargv = new char *[gsargc];
    for (int i = 0; i < gsargc; i++)
    {
        gsargv[i] = (char *)explodedCmd[i].c_str();
    }
    try
    {
        GhostscriptManager *gm = GhostscriptManager::GetInstance();
        gm->Execute(gsargc, gsargv);
        delete[] gsargv;
        res = 0;
    }
    catch (exception &e)
    {
        delete[] gsargv;
        msg << e.what();
        res = 1;
    }
}
```

**This code will be executed
on the Worker Thread**

```
void HandleOKCallback()
{
    Nan::HandleScope();
    Local<Value> argv[1];
    if (res == 0)
    {
        argv[0] = Null();
    }
    else
    {
        argv[0] = Error(Nan::New<String>(msg.str()).ToLocalChecked());
    }
    callback->Call(1, argv);
}
```

**When execute method end its
computation the results will be
reassembled and passed to
JS context**

```
private:
    string RAWcmd;
    int res;
    stringstream msg;
};
```

NAN_METHOD(Execute)

```
{
    Callback *callback = new Callback(info[1].As<Function>());
    Local<String> JScmd = Local<String>::Cast(info[0]);
    AsyncQueueWorker(new GhostscriptWorker(callback, *String::Utf8Value(JScmd)));
    info.GetReturnValue().SetUndefined();
}
```

Return immediately

Buffer example

```
void buffer_delete_callback(char* data, void* hint) {  
    free(data);  
}  
    Callback executed to free the memory allocated for the Buffer
```

Nicola Del Gobbo, 5 days ago

```
class GetKeyBufferWorker: public Nan::AsyncWorker {  
public:  
    GetKeyBufferWorker(Nan::Callback *callback, redis *db, std::string key)  
        :AsyncWorker(callback), db(db), key(key) {}  
    ~GetKeyBufferWorker() {}  
  
    void Execute() {  
        int rc;  
  
        std::stringstream strlenCmd;  
        strlenCmd << "STRLEN " + key;  
        rc = redis_exec(db, (strlenCmd.str()).c_str(), -1);  
        if(rc != REDIS_OK) {  
            // Handle error  
        }  
        /* Extract the return value of the last executed command (i.e. 'STRLEN test') " */  
        redis_value *strlen_result;  
        redis_exec_result(db, &strlen_result);  
        /* Cast the redis object to a string */  
        this->buffer_length = redis_value_to_int(strlen_result);  
  
        std::stringstream get_cmd;  
        get_cmd << "GET " + key;  
        rc = redis_exec(db, (get_cmd.str()).c_str(), -1);  
        if(rc != REDIS_OK) {  
            // Handle error  
        }  
        /* Extract the return value of the last executed command (i.e. 'GET test') " */  
        redis_value *get_result;  
        redis_exec_result(db, &get_result);  
        /* Cast the redis object to a string */  
        //this->buffer = static_cast<char*>(malloc(this->buffer_length));  
        this->buffer = const_cast<char*>(redis_value_to_string(get_result, 0));  
  
        GET command on database return a buffer  
  
        Buffers are Javascript objects, whose data is stored outside V8  
        anyways they are under V8's control  
  
        void HandleOKCallback() {  
            Nan::HandleScope();  
            int argc = 2;  
            Local<Value> argv[2];  
            argv[0] = Nan::Null();  
  
            // Not efficient solution to free memory  
            //argv[1] = Nan::CopyBuffer(this->buffer, this->buffer_length).ToLocalChecked();  
            Local<Object> data =  
                Nan::NewBuffer(this->buffer, this->buffer_length, buffer_delete_callback, this->buffer).ToLocalChecked();  
            argv[1] = data;  
            callback->Call(argc, argv);  
        }  
  
private:  
    redis *db;  
    std::string key;  
    char *buffer;  
    int buffer_length;  
};
```

```
NAN_PROPERTY_GETTER(Tensor::DataGetter) {  
    TensorflowNode::Tensor* tensor = ObjectWrap::Unwrap<TensorflowNode::Tensor>(info.This());  
    void* data = TF_TensorData(tensor->_tensor);  
    Sometimes the only way is to copy the buffer  
    size_t len = TF_TensorByteSize(tensor->_tensor);  
  
    // NewBuffer causes data to transfer internally, for example, when v8's GC collects this  
    // ArrayBuffer object, `data` would also be freed.  
    // The `data` is actually freed by tensorflow internally, so here CopyBuffer is the better way.  
    Local<Object> buffer = Nan::CopyBuffer((char*)data, len).ToLocalChecked();  
    info.GetReturnValue().Set(buffer);  
}
```

ObjectWrap API

- **ObjectWrap** is a way to expose your C++ code to JavaScript
- You have to extend **ObjectWrap** class that includes the plumbing to connect JavaScript code to a C++ object
- Classes extending **ObjectWrap** can be instantiated from JavaScript using the **new** operator, and their methods can be **directly invoked from JavaScript**
- Unfortunately, the **wrap** part really refers to a way to group methods and state, not decorate existing classes
- It's **your responsibility write custom code to bridge** each of your C++ class methods.

ObjectWrap example

```
NAN_MODULE_INIT(Database::Init) {
```

```
    Local<FunctionTemplate> tpl = Nan::New<FunctionTemplate>(New);  
    tpl->SetClassName(Nan::New("Database").ToLocalChecked());  
    tpl->InstanceTemplate()->SetInternalFieldCount(1);
```

Initialization code

Set prototype methods

```
Nan::SetPrototypeMethod(tpl, "getKey", GetKey);  
Nan::SetPrototypeMethod(tpl, "getKeyBuffer", GetKeyBuffer);  
Nan::SetPrototypeMethod(tpl, "getKeySync", GetKeySync);  
Nan::SetPrototypeMethod(tpl, "putKey", PutKey);  
Nan::SetPrototypeMethod(tpl, "putKeyBuffer", PutKeyBuffer);  
Nan::SetPrototypeMethod(tpl, "putKeySync", PutKeySync);
```

Set accessor methods

```
Local<ObjectTemplate> itpl = tpl->InstanceTemplate();  
Nan::SetAccessor(itpl, Nan::New("db_name").ToLocalChecked(), DbName);
```

Update constructor reference

```
constructor().Reset(v8::Isolate::GetCurrent(), Nan::GetFunction(tpl).ToLocalChecked());
```

```
Nan::Set(target, Nan::New("Database").ToLocalChecked(), Nan::GetFunction(tpl).ToLocalChecked());  
}
```

```
NAN_METHOD(Database::New) {
```

```
    // Here we need some control
```

```
    String::Utf8Value tmpDbName(info[0]->ToString());
```

```
    std::string dbName(*tmpDbName);
```

```
    if (info.IsConstructCall()) {
```

```
        KVDB::Database *database = new KVDB::Database(dbName);
```

```
        database->Wrap(info.This());
```

```
        info.GetReturnValue().Set(info.This());
```

```
    } else {
```

```
        const int argc = 1;
```

```
        Local<Value> argv[argc] = {info[0]};
```

```
        Local<Function> cons = Nan::New(constructor());
```

```
        info.GetReturnValue().Set(Nan::NewInstance(cons, argc, argv).ToLocalChecked());
```

```
    }
```

```
}
```

Use the new operator

Call as normal function

```
NAN_METHOD(Database::GetKey) {
```

```
    // Here we need some control
```

```
    String::Utf8Value tmpKey(info[0]->ToString());
```

```
    Nan::Callback *callback = new Nan::Callback(info[1].As<Function>());
```

```
    std::string key(*tmpKey);
```

```
    std::stringstream cmd;
```

```
    cmd << "GET " + key;
```

```
    KVDB::Database* database = ObjectWrap::Unwrap<KVDB::Database>(info.This());
```

```
    AsyncQueueWorker(new GetKeyWorker(callback, database->db, cmd.str()));
```

```
    info.GetReturnValue().SetUndefined();
```

```
}
```

Prototype method

```
NAN_GETTER(Database::DbName) {
```

```
    KVDB::Database* database = ObjectWrap::Unwrap<KVDB::Database>(info.This());
```

```
    info.GetReturnValue().Set(Nan::New(database->db_name).ToLocalChecked());
```

```
}
```

Accessor method

```
const mydb = new Database('test')
```

```
mydb.db_name
```

```
mydb.getKey('username', function(err, value) {
```

```
    if (err) {
```

```
        console.error(err)
```

```
    } else {
```

```
        console.log('Value from my async API')
```

```
        console.log(value)
```

```
    }
```

```
})
```




October 4-6, 2017 | Vancouver, BC

Lesson learned from my experience

- Create **simple interface** between C / C++ and JavaScript
- Insert **complexity** on the JavaScript side
- Use **ObjectWrap** API to return C / C++ object to JavaScript
- Validate the **types** on the native side
- Expose **only the functionalities you need**
- Don't block the event loop, stay **asynchronous**: this is a must if you will use the Addon on a distributed system
- Use **Buffer** to pass big quantity of data
- If you have many parameters use **JSON** to **pass** and **parse** data
- Consider the using of **protobuf**

Present and future


- The largest number of native addons is written using **NAN**
- **NAN** helps stay compatible with **old versions** of Node.js
- **N-API** will be a game changer on the native addon development
 - **ABI** compatibility
 - New **ES6** types
 - **Isolated** from V8 API
 - C / C++
 - Conversion tool that helps to migrate from **NAN**
 - Great **developer experience** for maintainers



LITTLE EXAMPLE

Implement high performance key value database

Binding to **Vedis** an embeddable datastore C library



```
const mydb = new Database('test')
```

```
mydb.db_name
```

```
mydb.putKeySync('username', 'NickNaso')
```

```
mydb.getKeySync('username')
```

```
mydb.putKey('password', 'keeplooking', function (err) {  
  if (err) {  
    console.error(err)  
  } else {  
    console.log("Value stored")  
  }  
})
```

```
mydb.getKey('username', function(err, value) {  
  if (err) {  
    console.error(err)  
  } else {  
    console.log('Value from my async API')  
    console.log(value)  
  }  
})
```

```
const buffer = Buffer.from('qwertyuiopasdfghjklzxcvbnm1234567890')
```

```
mydb.putKeyBuffer('image', buffer, function (err) {  
  if (err) {  
    console.error(err)  
  } else {  
    console.log('Buffer stored')  
  }  
})
```

```
mydb.getKeyBuffer('image', function (err, buffer) {  
  if (err) {  
    console.error(err)  
  } else {  
    console.log(buffer.toString())  
  }  
})
```

Thanks!



nicoladelgobbo@gmail.com

[@NickNaso](#) on Twitter

All examples and materials are here

<https://github.com/NickNaso/conf-ni-2017>